

JUDCon

JBoss Users & Developers Conference

Boston:2011

DEMISTIFYING COMPLEX EVENT PROCESSING

Edson Tirelli

Drools Fusion Lead, Red Hat

May 2nd, 2011

Agenda

- Brief introduction on CEP and Terminology
- Drools Fusion: Complex Event Processing extensions
 - Event Declaration and Semantics
 - Event Cloud, Streams and the Session Clock
 - Temporal Reasoning
 - Sliding Window Support
 - Streams Support
 - Memory Management
- Questions & Answers

Terminology: Event

“An event is an **observable occurrence**.”

“An event in the Unified Modeling Language is a **notable occurrence** at a particular **point in time**.”

<http://www.wikipedia.org>

“Anything that **happens**, or is **contemplated as happening**.”

“An **object that represents, encodes or records an event**, generally for the purpose of **computer processing**”

<http://complexevents.com>

Terminology: Event

For the scope of this presentation:

“An event is a **significant change of state** at a particular **point in time**”

Terminology: Complex Event

“**Complex Event**, is an abstraction of other events called its members.”

- **Examples:**

- **The 1929 stock market crash** – an abstraction denoting many thousands of member events, including individual stock trades)
- **The 2004 Indonesian Tsunami** – an abstraction of many natural events
- **A completed stock purchase** -an abstraction of the events in a transaction to purchase the stock
- **A successful on-line shopping cart checkout** – an abstraction of shopping cart events on an on-line website

- *Source: <http://complexevents.com>*

Terminology: CEP

“**Complex Event Processing**, or CEP, is primarily an event processing concept that deals with the task of processing multiple events with the goal of **identifying the meaningful events** within the event cloud.

CEP employs techniques such as **detection** of complex patterns of many events, event **correlation** and **abstraction**, event hierarchies, and relationships between events such as causality, membership, and timing, and event-driven processes.”

-- wikipedia

Terminology: CEP

- Examples:

- Emergency Response Systems
- Credit Card Fraud Detection
- Logistics Real-Time Awareness solution
- Neonatal ICU: infant vital signs monitoring

Terminology: CEP vs ESP

Complex Event Processing, or CEP, and **Event Stream Processing**, or ESP, are two technologies that were born separate, but **converged**.

- *An oversimplification:* In their origins...
 - **Event Stream Processing** focused on the ability to process high volume **streams** of events.
 - **Complex Event Processing** focused on defining, detecting and processing the **relationships** among events.

Terminology: CEP and ESP

For the scope of this presentation:

“**CEP** is used as a common term meaning both **CEP** and **ESP**.”

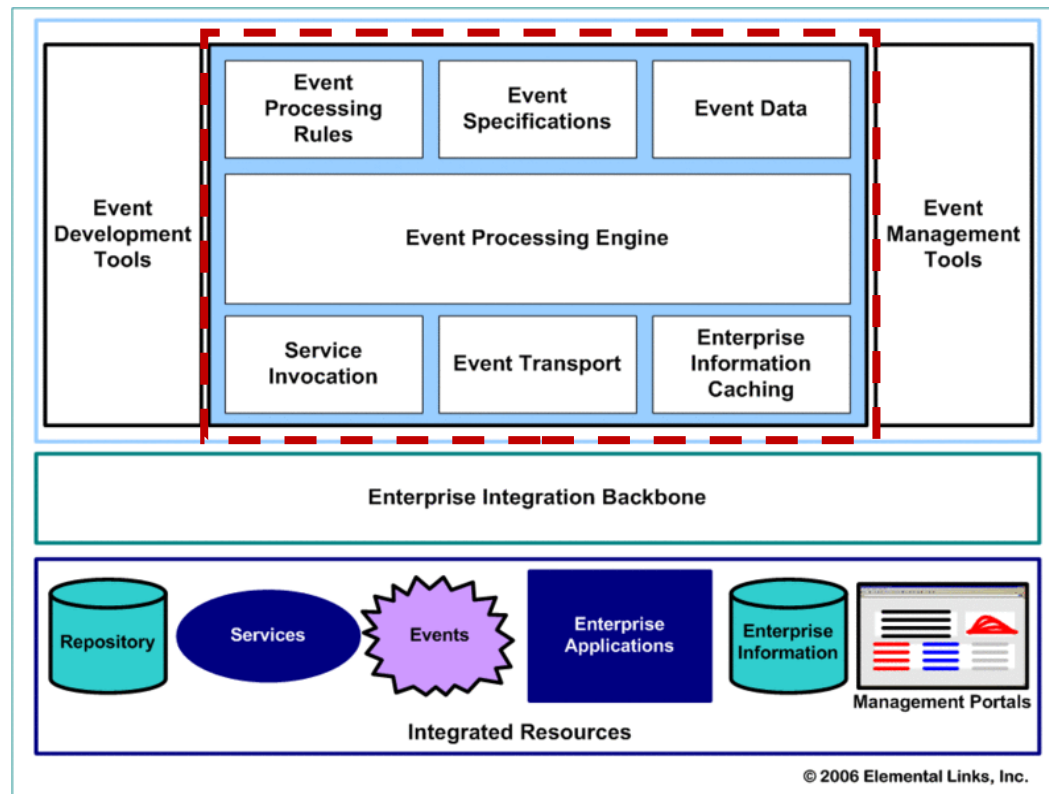
Terminology: EDA

“**Event Driven Architecture (EDA)** is a software architecture pattern promoting the **production**, **detection**, **consumption** of, and **reaction** to events. An **event** can be defined as "a significant change in state"[1]. For example, when a consumer purchases a car, the car's state changes from "for sale" to "sold". A car dealer's system architecture may treat this state change as an event to be produced, published, detected and consumed by various applications within the architecture.”

http://en.wikipedia.org/wiki/Event_Driven_Architecture

EDA vs CEP

CEP is a **component** of the EDA

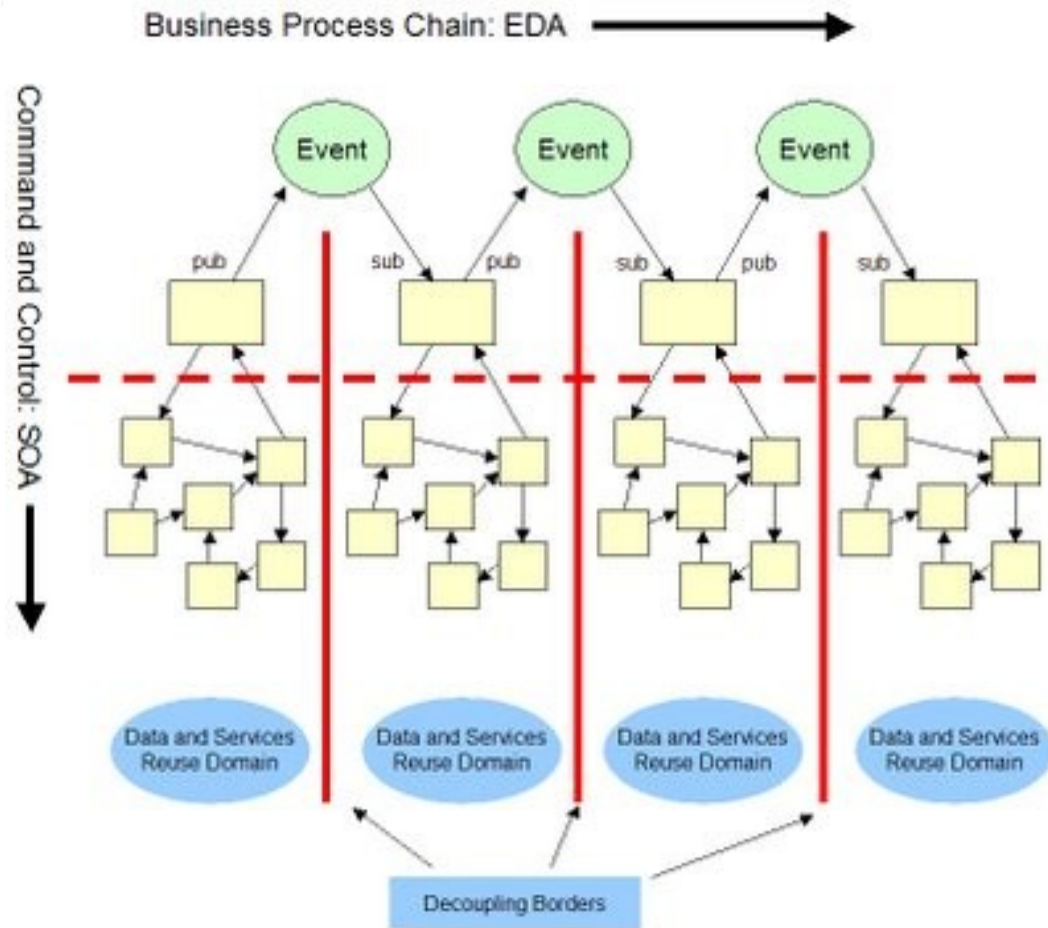


Source: http://elementallinks.typepad.com/.shared/image.html?/photos/uncategorized/simple_event_flow.gif

EDA vs SOA

- EDA is ****not**** SOA 2.0
- Complementary architectures
- Metaphor
 - In our **body**:
 - SOA is used to build our **muscles and organs**
 - EDA is used to build our **sensory system**

EDA vs SOA



Source: <http://soa-eda.blogspot.com/2006/11/how-eda-extends-soa-and-why-it-is.html>

Complex Event Processing

- **A few characteristics of common CEP scenarios:**
 - Huge volume of events, but only a few of real interest
 - Usually events are immutable
 - Usually queries/rules have to run in reactive mode
 - Strong temporal relationships between events
 - Individual events are usually not important
 - The composition and aggregation of events is important

Drools Vision



Guvnor 

Expert 

Fusion 

Flow 

“A **common platform** to **model** and **govern** the business **logic** of the enterprise.”



JBoss Enterprise BRMS v5.2

Drools Fusion: Enables...

- **Event Detection:**
 - From an event cloud or set of streams, select all the meaningful events, and only them.
- **[Temporal] Event Correlation:**
 - Ability to correlate events and facts declaring both temporal and non-temporal constraints between them.
 - Ability to reason over event aggregation
- **Event Abstraction:**
 - Ability to compose complex events from atomic events AND reason over them

Drools Fusion

◦ Features:

- Event Semantics as First Class Citizens
- Allow Detection, Correlation and Composition
- Temporal Constraints
- Session Clock
- Stream Processing
- Sliding Windows
- CEP volumes (scalability)
- (Re)Active Rules
- Data Loaders for Input

Demo

- **Twitter Stream CEP Demo:**
 - Listen to the Twitter Stream API
 - Twitter4J API
 - Listens to a random sample of tweets
 - Detects patterns and reacts
 - Drools Fusion
 - Simple one process (multi-thread) demo
 - Focus on specific features

Event Declaration and Semantics

```
// declaring existing class
import some.package.VoiceCall
declare VoiceCall
  @role( event )
  @timestamp( calltime )
  @duration( duration )
end

// generating an event class
declare StockTick
  @role( event )

  symbol : String
  price : double
end
```

Event semantics:

- Point-in-time and Interval

An event is a fact with a few special characteristics:

- Usually immutable, but not enforced
- Strong temporal relationships
- Lifecycle may be managed
- Allow use of sliding windows

“All events are facts, but not all facts are events.”

Temporal Reasoning

- Semantics for:
 - **time:** discrete
 - **events:** point-in-time and interval
 - Ability to express temporal relationships:
 - Allen's 13 temporal operators
-
- **James F. Allen** defined the 13 possible temporal relations between two events.
 - **Eiko Yoneki** and **Jean Bacon** defined a unified semantics for event correlation over time and space.

Temporal Relationships

rule “Shipment not picked up in time”

when

Shipment(\$pickupTime : scheduledPickupTime)

not ShipmentPickup(this before \$pickupTime)

then

// shipment not picked up... action required.

end

Temporal Relationships

rule “Shipment not picked up in time”

when

Shipment(\$pickupTime : scheduledPickupTime)

not ShipmentPickup(this before \$pickupTime)

then

// shipment not picked up... Action required.















end




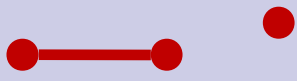
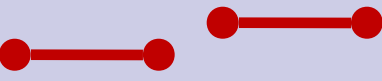


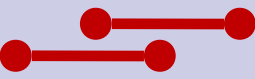



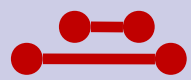


Temporal
Relationship

The diagram consists of a black rectangular box at the bottom containing the text 'Temporal Relationship'. A vertical black line extends upwards from the top center of this box, ending in an arrowhead that points to the underlined text 'this' in the rule's condition line: 'not ShipmentPickup(this before \$pickupTime)'.

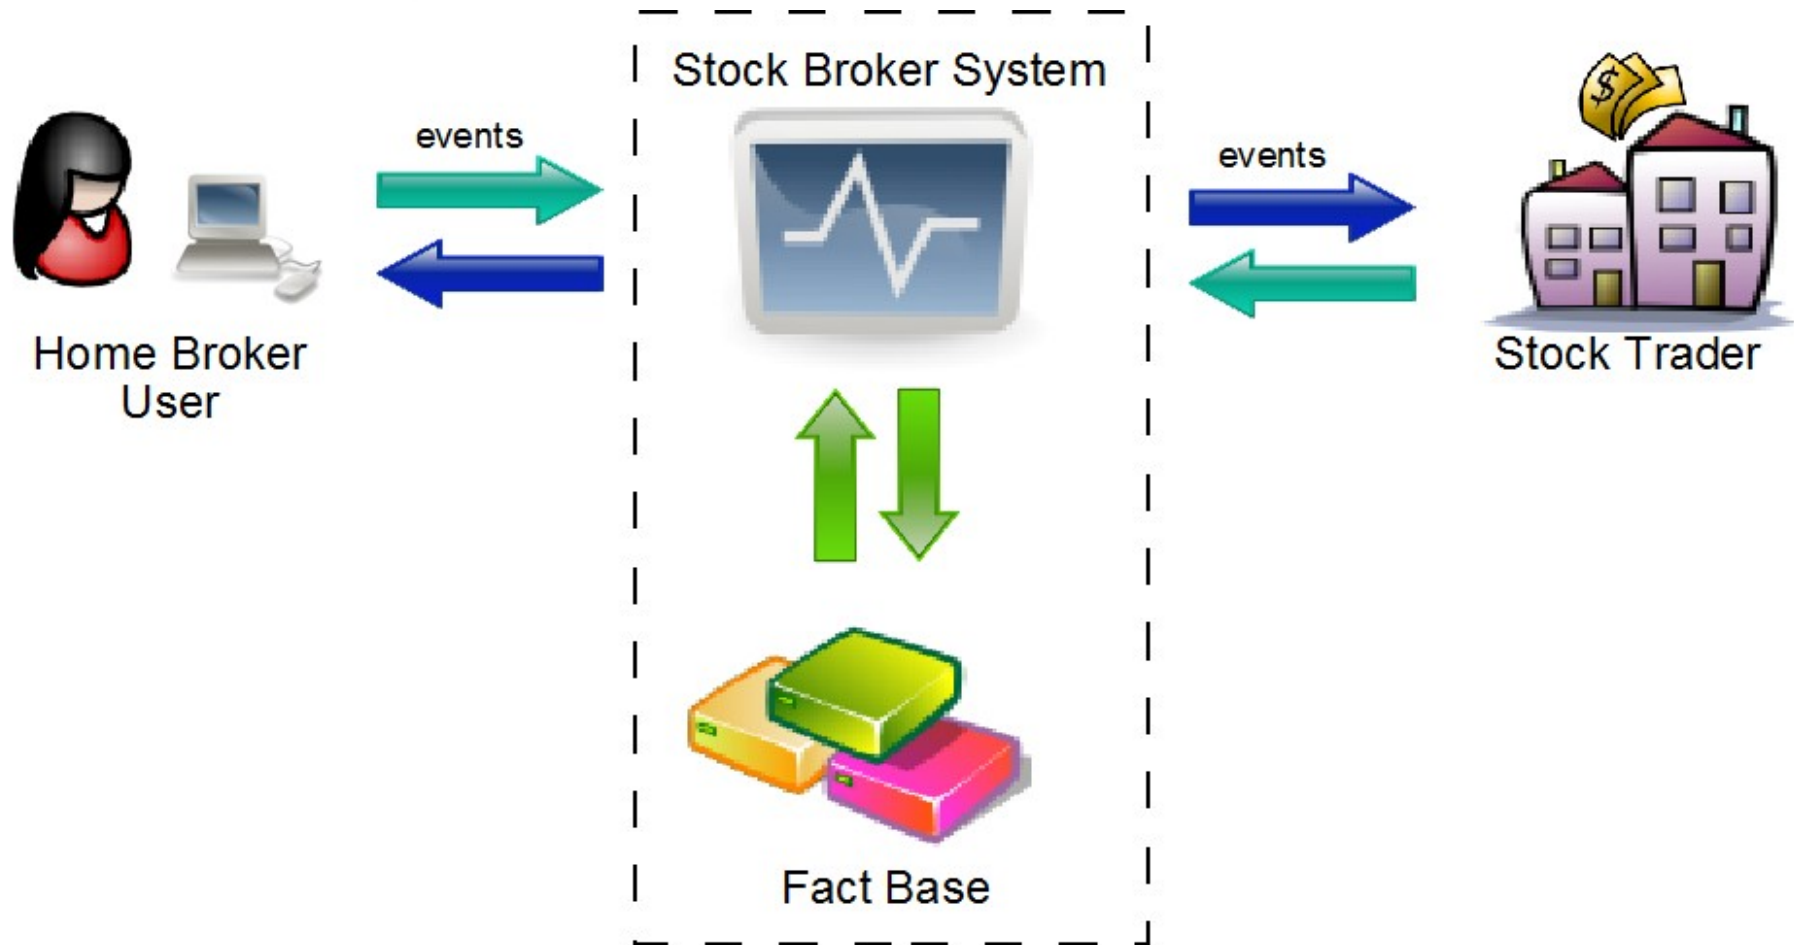
Allen's 13 Temporal Operators

		Point-Point	Point-Interval	Interval-Interval
A before B	<i>A</i> <i>B</i>			
A meets B	<i>A</i> <i>B</i>			
A overlaps B	<i>A</i> <i>B</i>			
A finishes B	<i>A</i> <i>B</i>			
A includes B	<i>A</i> <i>B</i>			
A starts B	<i>A</i> <i>B</i>			
A coincides B	<i>A</i> <i>B</i>			

Allen's 13 Temporal Operators

		Point-Point	Point-Interval	Interval-Interval
A after B	$\begin{matrix} A \\ B \end{matrix}$			
A metBy B	$\begin{matrix} A \\ B \end{matrix}$			
A overlapedBy B	$\begin{matrix} A \\ B \end{matrix}$			
A finishedBy B	$\begin{matrix} A \\ B \end{matrix}$			
A during B	$\begin{matrix} A \\ B \end{matrix}$			
A finishes B	$\begin{matrix} A \\ B \end{matrix}$			

Streams : Simple Example Scenario



Stream Support (entry-points)

- **A scoping abstraction for stream support**
 - Rule compiler gather all entry-point declarations and expose them through the session API
 - Engine manages all the scoping and synchronization behind the scenes.

```
rule "Stock Trade Correlation"
when
    $c : Customer( type == "VIP" )
    BuyOrderEvent( customer == $c, $id : id ) from entry-point "Home Broker Stream"
    BuyAckEvent( sourceEvent == $id ) from entry-point "Stock Trader Stream"
then
    // take some action
end
```


Cloud Mode, Stream Mode, Session Clock

CLOUD

- No notion of “flow of time”: the engine sees all facts without regard to time
- No attached Session Clock
- No requirements on event ordering
- No automatic event lifecycle management
- No sliding window support

STREAM

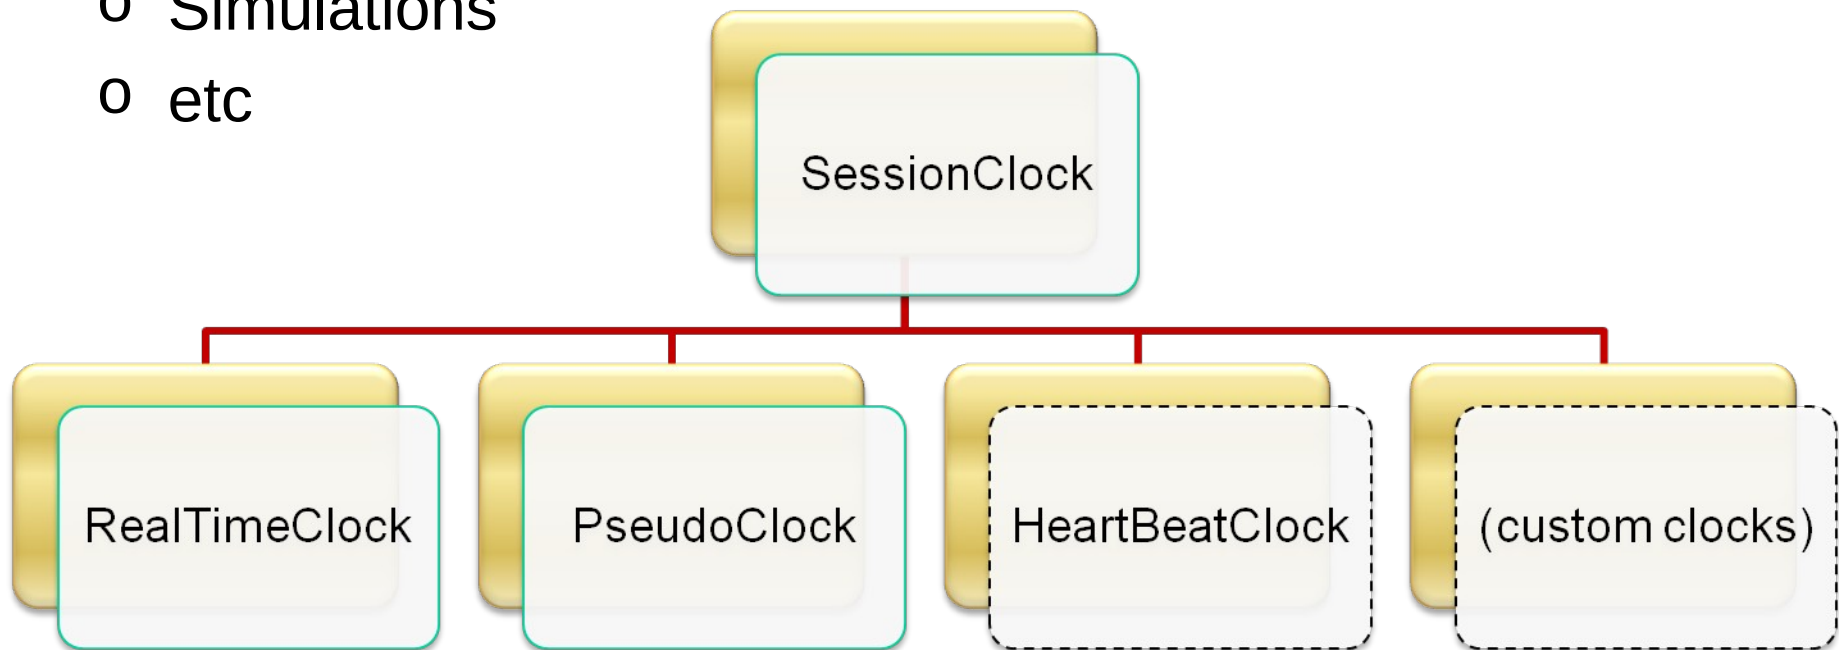
- Notion of “flow of time”: concept of “now”
- Session Clock has an active role synchronizing the reasoning
- Event Streams must be ordered
- Automatic event lifecycle management
- Sliding window support
- Automatic rule delaying on absence of facts

Reference Clock

- Reference clock defines the flow of time
- **Named Session Clock**
 - is assigned to each session created
- **Synchronizes time sensitive operations**
 - duration rules
 - event streams
 - process timers
 - sliding windows

Session Clock

- 0 Uses the strategy pattern and multiple implementations:
 - 0 Real-time operation
 - 0 Tests
 - 0 Simulations
 - 0 etc



Session Clock

- Selecting the session clock:

- API:


```
KnowledgeSessionConfiguration conf = ...  
conf.setOption( ClockTypeOption.get( "realtime" ) );
```

- System Property or Configuration File:

```
drools.clockType = pseudo
```


Sliding Window Support

- Allows reasoning over a moving window of “interest”
 - Time
 - Length



Sliding Window Support

- Allows reasoning over a moving window of “interest”
 - Time
 - Length



Delaying Rules

- 0 Negative patterns may require rule firings to be delayed.



```
rule "Order timeout"  
when  
    $bse : BuyShares ( $id : id )  
    not BuySharesAck( id == $id, this after[0s,30s] $bse )  
then  
    // Buy order was not acknowledged. Cancel operation  
    // by timeout.  
end
```

Temporal Dimension

- Requires the support to the temporal dimension
 - A rule/query might match in a given point in time, and not match in the subsequent point in time
- That is the single most difficult requirement to support in a way that the engine:
 - stays deterministic
 - stays a high-performance engine
- Achieved mostly by compile time optimizations that enable:
 - constraint tightening
 - match space narrowing
 - memory management

Q&A

- **Drools project site:**
 - <http://www.drools.org> (<http://www.jboss.org/drools/>)
- **Documentation:**
 - <http://www.jboss.org/drools/documentation.html>



Edson Tirelli – etirelli@redhat.com

JUDCon

JBoss Users & Developers Conference

Boston:2011